

Introduction to Functional programming

David Li

Functions

Functions in R

- A core activity of an R programmer.
- “user” → developer
- When to write a function
 - Encapsulate a sequence of expressions that need to be executed numerous times, perhaps under slightly different conditions.
 - Code must be shared with others or the public
- Create an interface to the code: via a set of parameters.
- This interface provides an abstraction of the code to potential users.
 - Ex: `sort()`

Your First Function

```
f <- function() {  
  ## This is an empty function  
}  
## Functions have their own class  
class(f)  
# Execute this function  
f()
```

```
#more fun  
f <- function() {  
  cat("Hello, world!\n")  
}  
f()
```

```
#with a parameter  
f <- function(num) {  
  for(i in seq_len(num)) {  
    cat("Hello, world!\n")  
  }  
}  
f(3)
```

```
# with return value  
f <- function(num) {  
  hello <- "Hello, world!\n"  
  for(i in seq_len(num)) {  
    cat(hello)  
  }  
  chars <- nchar(hello) * num  
  chars # logical last expression  
  returned  
}  
meaningoflife = f(3)
```

#return the very last expression that is evaluated.

Default value

```
f()
f <- function(num = 1) {
  hello <- "Hello, world!\n"
  for(i in seq_len(num)) {
    cat(hello)
  }
  chars <- nchar(hello) * num
  chars
}
f() ## Use default value for 'num'
```

```
f(2)
f(num=2) #specified using argument name
```

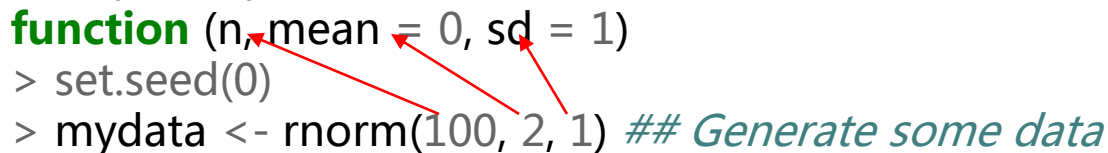
So far, we have written a function that

- has one *formal argument* named `num` with a *default value* of 1.
- prints the message "Hello, world!" to the console a number of times indicated by the argument `num`
- *returns* the number of characters printed to the console

Argument Matching

- R functions arguments can be matched *positionally* or by name.
- Positional matching just means that R assigns the first value to the first argument, the second value to second argument, etc.

```
> str(rnorm)
function (n, mean = 0, sd = 1)
> set.seed(0)
> mydata <- rnorm(100, 2, 1) ## Generate some data
```



100 is assigned to the n argument, 2 is assigned to the mean argument, and 1 is assigned to the sd argument, all by positional matching.

Specifying arguments by name

- Order doesn't matter then

```
> sd(na.rm = FALSE, mydata)
```

Here, the mydata object is assigned to the x argument, because it's the only argument not yet specified.

- Function arguments can also be *partially* matched
- The order of operations when given an argument
 1. Check for exact match for a named argument
 2. Check for a partial match
 3. Check for a positional match

Example

```
> args(lm)
function (formula, data, subset, weights, na.action, method = "qr",
model = TRUE, x = FALSE, y = FALSE, qr = TRUE, singular.ok = TRUE,
contrasts = NULL, offset, ...)
NULL
```

The following two calls are equivalent.

```
set.seed(0)
mydata = data.frame(y=rnorm(20), x=rnorm(20))
lm(data = mydata, y ~ x, model = FALSE, 1:20)
lm(y ~ x, mydata, 1:20, model = FALSE)
```


Lazy Evaluation

- Arguments to functions are evaluated *lazily*, so they are evaluated only as needed in the body of the function.

```
> f <- function(a, b) {  
  a^2  
}  
> f(2)
```

```
> f <- function(a, b) {  
  print(a)  
  print(b)  
}  
> f(45)
```

The ... Argument

- A special argument in R
- Indicate a variable number of arguments that are usually passed on to other functions.

```
myplot <- function(x, y, type = "|", ...) {  
  plot(x, y, type = type, ...) ## Pass '...' to 'plot' function  
}
```

- The ... argument is necessary when the number of arguments passed to the function cannot be known in advance.

```
> args(paste)  
function (... , sep = " ", collapse = NULL)  
NULL  
> args(cat)  
function (... , file = "", sep = " ", fill = FALSE, labels = NULL, append = FALSE)  
NULL
```

Arguments Coming After the ... Argument

- One catch with ... is that any arguments that appear *after* ... on the argument list must be named explicitly and cannot be partially matched or matched positionally.

> **args**(**paste**)

function (... , sep = " ", collapse = **NULL**)
NULL

paste("a", "b", sep = "+")

paste("a", "b", se = "+")

Summary

- Functions can be defined using the `function()` directive and are assigned to R objects just like any other R object
- Functions can be defined with named arguments; these function arguments can have default values
- Functions arguments can be specified by name or by position in the argument list
- Functions always return the last expression evaluated in the function body
- A variable number of arguments can be specified using the special `...` argument in a function definition.

Example: Newton-Rapson method to find a square root of an integer number

It is an iterative number method. This type of methods create a succession $(x_0, x_1 \cdots x_n)$.

With some initial conditions given a root of the function f , α

$$f(\alpha) = 0$$

$$\lim_{n \rightarrow \infty} \alpha - x_n = 0$$

More exactly newton-rapson follow the next schema:

$$x_{n+1} = x_n - \frac{f(x)}{f'(x_n)}$$


For finding a square root of an integer (t) you can do this trick.

$$f(x) = x^2 - t$$

then in our case:

$$x_{n+1} = x_n - \frac{(x_n)^2 - t}{2x_n}$$

where $f'(x) = 2x$.



Functional Programming

Looping on the Command Line

- `apply()`: Apply a function over the margins of an array
- `lapply()`: Loop over a list and evaluate a function on each element
- `sapply()`: Same as `lapply` but try to simplify the result
- `mapply()`: Multivariate version of `lapply`
- `tapply()`: Apply a function over subsets of a vector

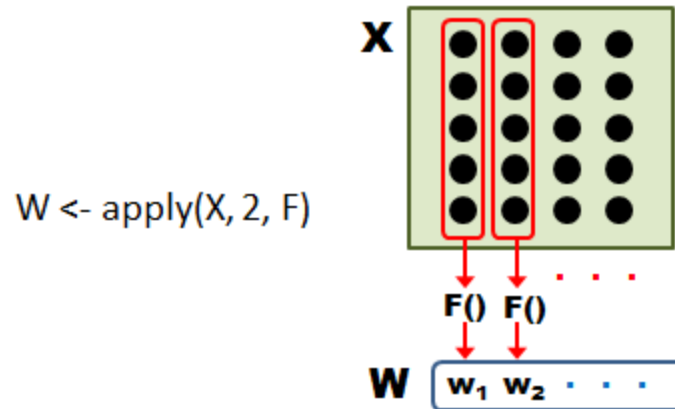

apply()

- Evaluate a function (often an anonymous one) over the margins of an array.
- Most often, apply a function to the rows or columns of a matrix (which is just a 2-dimensional array). Also, applicable to general arrays

```
> str(apply)  
function (X, MARGIN, FUN, ...)
```

- The arguments to apply() are
 - X is an array (matrix is just a 2D array)
 - MARGIN is an integer vector indicating which margins should be “retained” .
 - FUN is a function to be applied
 - ... is for other arguments to be passed to FUN

apply()



apply()

- Examples

```
> set.seed(0)
> x <- matrix(rnorm(200), 20, 10)
> apply(x, 2, mean) ## Take the mean of each column
> apply(x, 1, sum)  ## Take the mean of each row
> a <- array(rnorm(2 * 2 * 10), c(2, 2, 10))
> apply(a, c(1, 2), mean)
```

- Shortcuts

- rowSums = apply(x, 1, sum)
- rowMeans = apply(x, 1, mean)
- colSums = apply(x, 2, sum)
- colMeans = apply(x, 2, mean)

lapply()

- The lapply() function does the following simple series of operations:
 1. it loops over a list, iterating over each element in that list
 2. it applies a *function* to each element of the list (a function that you specify)
 3. and returns a list (the l is for “list”).
- This function takes three arguments
 - (1) a list X, If X is not a list, **it will be coerced to a list using as.list()**.
 - (2) a function (or the name of a function) FUN;
 - (3) other arguments via its ... argument.

lapply()

- the actual looping is done internally in C code for efficiency reasons.

```
> lapply
function (X, FUN, ...)
{
  FUN <- match.fun(FUN)
  if (!is.vector(X) || is.object(X))
  X <- as.list(X)
  .Internal(lapply(X, FUN))
}
<bytecode: 0x7fcc8388f758>
<environment:
namespace:base>
```

lapply()

- Example 1

```
> x <- list(a = 1:5, b = rnorm(10))  
> lapply(x, mean)  
> x <- list(a = 1:4, b = rnorm(10), c = rnorm(20, 1), d = rnorm(100, 5))  
> lapply(x, mean)
```

- Example 2

```
> x <- 1:4  
> lapply(x, runif)
```

When you pass a function to `lapply()`, `lapply()` takes elements of the list and passes them as the *first argument* of the function you are applying.

lapply()

- the ... Argument

```
> x <- 1:4  
> lapply(x, runif, min = 0, max = 10)
```

- *anonymous* functions.


```
> x <- list(a = matrix(1:4, 2, 2), b = matrix(1:6, 3, 2))  
> lapply(x, function(elt) { elt[,1] })
```

V.S.

```
> f <- function(elt) {  
  elt[, 1]  
}  
> lapply(x, f)
```

sapply()

- The sapply() function behaves similarly to lapply(); the only real difference is in the return value.



- sapply() will try to simplify the result of lapply() if possible. Essentially, sapply() calls lapply() on its input and then applies the following algorithm
 - If the result is a list where every element is length 1, then a vector is returned
 - If the result is a list where every element is a vector of the same length (> 1), a matrix is returned.
 - If it can't figure things out, a list is returned
- ```
> x <- list(a = 1:4, b = rnorm(10), c = rnorm(20, 1), d = rnorm(100, 5))
> lapply(x, mean)
> sapply(x, mean)
```


# split()

- The combination of split() and a function like lapply() or sapply() is a common paradigm in R.
  - > **library**(datasets)
  - > **head**(airquality)
  - > s <- **split**(airquality, airquality\$Month)
  - > str(s)
  - > **lapply**(s, **function**(x) {  
    **colMeans**(x[, **c**("Ozone", "Solar.R", "Wind")])  
})
  - > **sapply**(s, **function**(x) {  
    **colMeans**(x[, **c**("Ozone", "Solar.R", "Wind")])  
})
  - > **sapply**(s, **function**(x) {  
    **colMeans**(x[, **c**("Ozone", "Solar.R", "Wind")], na.rm = **TRUE**)  
})



# tapply

- `tapply()` is used to apply a function over subsets of a vector. It can be thought of as a combination of `split()` and `apply()` for vectors only.



> str(**tapply**)

**function** (X, INDEX, FUN = **NULL**, ..., simplify = **TRUE**)

- The arguments to `tapply()` are as follows:
  - X is a vector
  - INDEX is a factor or a list of factors (or else they are coerced to factors)
  - FUN is a function to be applied
  - ... contains other arguments to be passed FUN
  - simplify, should we simplify the result?


# tapply

```
> ## Simulate some data
> x <- c(rnorm(10), runif(10), rnorm(10, 1))
> ## Define some groups with a factor variable
> f <- gl(3, 10)
> f
[1] 1 1 1 1 1 1 1 1 1 1 2 2 2 2 2 2 2 2 2 2 3 3 3 3 3 3 3 3 3 3
Levels: 1 2 3
> tapply(x, f, mean)

> tapply(x, f, mean, simplify = FALSE)
> tapply(x, f, mean, simp = FALSE) #will it return the same? why
when returning >1 value. tapply() will not simplify the result and
will return a list.
> tapply(x, f, range)
```

# mapply()

- A multivariate apply of sorts which applies a function in parallel over a set of arguments.
- Recall that lapply() and friends only iterate over a single R object. What if you want to iterate over multiple R objects in parallel? This is what mapply() is for.



```
W <- mapply(F, X, Y, ...)
```

```
> str(mapply)
```

**function** (FUN, ..., MoreArgs = **NULL**, SIMPLIFY = **TRUE**, USE.NAMES = **TRUE**)

The arguments to mapply() are

- FUN is a function to apply
- ... contains R objects to apply over
- MoreArgs is a list of other arguments to FUN.
- SIMPLIFY indicates whether the result should be simplified

# mapply()

```
list(rep(1, 4), rep(2, 3), rep(3, 2), rep(4, 1))
```

```
> mapply(rep, 1:4, 4:1)
```

```
> noise <- function(n, mean, sd) {
 rnorm(n, mean, sd)
}
```

```
> ## Simulate 5 random numbers
```

```
> noise(5, 1, 2)
```

```
[1] -0.5196913 3.2979182 -0.6849525 1.7828267 2.7827545
```

```
>
```

```
> ## This only simulates 1 set of numbers, not 5
```

```
> noise(1:5, 1:5, 2)
```

```
[1] -1.670517 2.796247 2.776826 5.351488 3.422804
```

```
> mapply(noise, 1:5, 1:5, 2) == list(noise(1, 1, 2), noise(2, 2, 2),
 noise(3, 3, 2), noise(4, 4, 2),
 noise(5, 5, 2))
```

# Vectorizing a Function

- The `mapply()` function can be used to automatically “vectorize” a function: take a function that typically only takes single arguments and create a new function that can take vector arguments.

```
> sumsq <- function(mu, sigma, x) {
 sum(((x - mu) / sigma)^2)
}

> x <- rnorm(100) ## Generate some data
> sumsq(1:10, 1:10, x) ## This is not what we want
[1] 110.2594
```

However, we can do what we want to do by using `mapply()`.

```
> mapply(sumsq, 1:10, 1:10, MoreArgs = list(x = x))
```

# Vectorize()

- It can automatically create a vectorized version of your function.
- Example: create a vsumsq() function that is fully vectorized as follows.

```
> vsumsq <- Vectorize(sumsq, c("mu", "sigma"))
> vsumsq(1:10, 1:10, x)
[1] 196.2289 121.4765 108.3981 104.0788 102.1975 101.2393 100.6998
[8] 100.3745 100.1685 100.0332
```

# Summary

- The loop functions in R are very powerful because they allow you to conduct a series of operations on data using a compact form
- The operation of a loop function involves iterating over an R object (e.g. a list or vector or matrix), applying a function to each element of the object, and then collating the results and returning the collated results.
- Loop functions make heavy use of anonymous functions, which exist for the life of the loop function but are not stored anywhere
- The `split()` function can be used to divide an R object into subsets determined by another variable which can subsequently be looped over using loop functions.

# Computing Lab Ex

- Further Readings
  - <http://adv-r.had.co.nz/Functionals.html>
  - <https://towardsdatascience.com/functional-programming-in-r-with-purrr-469e597d0229>
- Computing Lab Ex
  - Try to run the sample code in notebook
  - Speed up the loop operation in R
    - <http://stackoverflow.com/questions/2908822/speed-up-the-loop-operation-in-r>
  - Is R's apply family more than syntactic sugar?
    - <http://stackoverflow.com/questions/2275896/is-rs-apply-family-more-than-syntactic-sugar>